

CPSC 4330: Big Data Analytics

Writing a MapReduce Program in
Java

A Sample MapReduce Program: WordCount

- The program consists of three files
 - The Mapper
 - The Reducer
 - The driver code

<https://hadoop.apache.org/docs/r3.4.1/api/>

Keys and Values are Objects

- Keys and values in Hadoop are Java Objects
 - Not Java primitives. Cannot use things like int or float as keys or values
 - Should use the corresponding Writable subclass (e.g. IntWritable, FloatWritable) to “wrap” the values.
- Values are objects which implement *Writable*
 - The writable interface is used in Hadoop’s serialization process, allowing these values to be efficiently stored to or read from files, as well as being passed across the network.
- Keys are objects which implement *WritableComparable*
 - Why?
 - To make it possible to sort the keys

What is Writable?

- The *Writable* interface makes serialization quick and easy for Hadoop
- Any value's type must implement the *Writable* interface
- Hadoop defines its own 'box classes' for strings, integers, and so on
 - IntWritable for ints
 - LongWritable for longs
 - FloatWritable for floats
 - DoubleWritable for doubles
 - Text for strings
 - etc.

```
e.g.  
IntWritable key = new IntWritable(5);  
  
int number = key.get();
```

What is WritableComparable?

- A WritableComparable is a Writable which is also Comparable
 - Two WritableComparables can be compared against each other to determine their 'order'
 - Keys must be WritableComparables because they are passed to the Reducer in sorted order
- Note that despite their names, all Hadoop box classes implement both Writable and WritableComparable
 - For example, IntWritable is actually a WritableComparable

WordCount Program – The Mapper

Input Data (HDFS file)

```
the cat sat on the mat
the aardvark sat on the sofa
...
```

Record Reader

Key	Value
0	the cat sat on the mat
23	the aardvark sat on the sofa
52	...
...	...

Mapper

map()

map()

key	value
the	1
cat	1
sat	1
on	1
the	1
mat	1

key	value
the	1
aardvark	1
sat	1
on	1
the	1
sofa	1

The Mapper: Complete Code

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String line = value.toString();

        for (String word : line.split("\\W+")) {
            if (word.length() > 0) {
                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

The Mapper: import Statements

```
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;
```

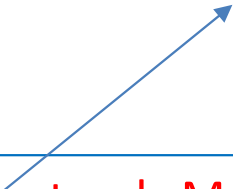
You will typically import the two classes `IOException` and `Mapper` (highlighted in red), in every Mapper program you write. Other classes such as `IntWritable`, `LongWritable`, `Text` are imported if you need to use those `WritableComparable` in your program.

The Mapper: Main Code

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        String line = value.toString();  
  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}
```

The Mapper: Class Declaration (1)

The WordMapper class extends the base class *Mapper*



```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        String line = value.toString();  
  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}
```

The Mapper: Class Declaration (2)

Input key and value types

Intermediate output key and value types

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        String line = value.toString();  
  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}
```

- Keys must be WritableComparable; Values must be Writable
- Type of Input key and value are determined by whatever InputFormat you specify. Can specify the InputFormat in Driver class (discuss later). Default is TextInputFormat.
- Programmers choose what the output key and value types are.

The Mapper: The map method

Types of input key and value specified in the class signature and the map method's signature must match

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        String line = value.toString();  
  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}
```

Context object is used to write intermediate data. It also contains information about the job's configuration.

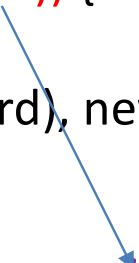
The map Method: Processing the line (1)

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        String line = value.toString();  
  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}
```

—————→ A text object is not a string. toString() method returns the string that a text object contains.

The map Method: Processing the line (2)

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        String line = value.toString();  
  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}
```




- Split the string up into words using the specified delimiter and then loop through the words
- `\W`: Matches any non-word character.
- `+`: Matches the previous element one or more times.
- The backslash “\” before “\W+” is an escape character, indicates that the character that follows it either is a special character, or should be interpreted literally.

Word character:
a-z, A-Z, 0-9, _

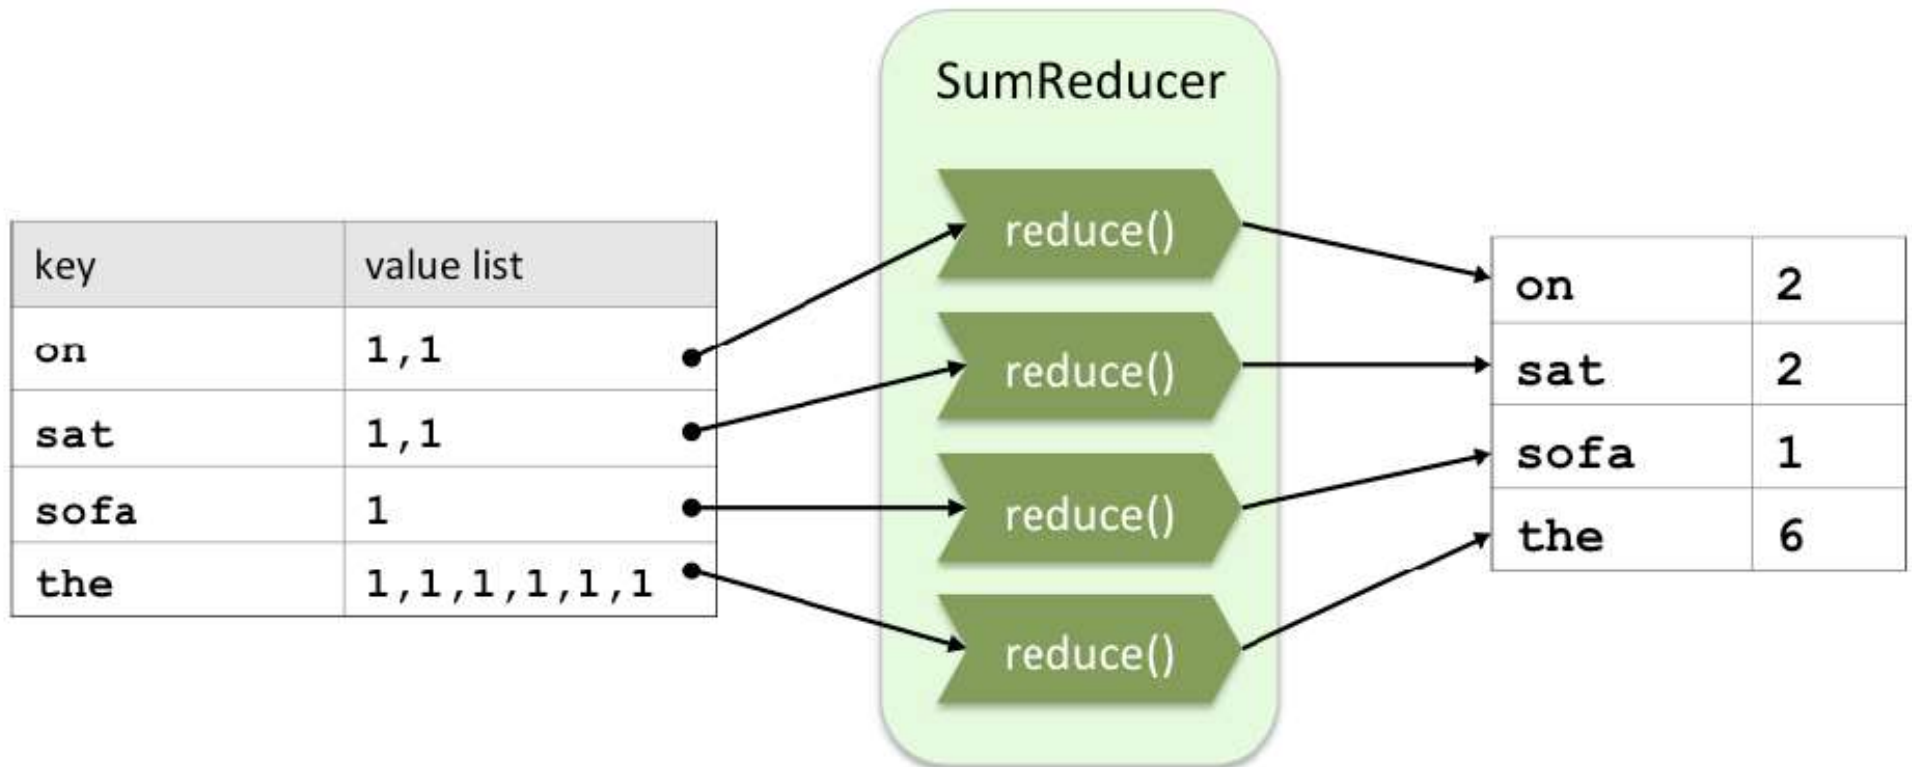
The map Method: Outputting Intermediate Data

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        String line = value.toString();  
  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}
```



- Output a (key, value) pair
- Key: word
- Value: 1
- No I/O code: Hadoop takes care of I/O

WordCount Program – The Reducer



The Reducer: Complete Code

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int wordCount = 0;

        for (IntWritable value : values) {
            wordCount += value.get();
        }

        context.write(key, new IntWritable(wordCount));
    }
}
```

The Reducer: Import Statements

```
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;
```

You will typically import *java.io.IOException*, and the *org.apache.hadoop.mapreduce.Reducer* classes shown, in every Reducer you write. Other classes like *IntWritable*, *Text* are imported if your program needs to use those *WritableComparable*.

The Reducer: Main Code

SumReducer class extends the
Reducer base class

Intermediate key
and value types


Final output key
and value types

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
  
        int wordCount = 0;  
  
        for (IntWritable value : values) {  
            wordCount += value.get();  
        }  
  
        context.write(key, new IntWritable(wordCount));  
    }  
}
```

The Reducer: reduce method

- The reduce method receives a key and an Iterable collection of objects (which are values output from the Mappers for that key)
- Types of key and value must match the types of input key value in the class signature

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
  
        int wordCount = 0;  
  
        for (IntWritable value : values) {  
            wordCount += value.get();  
        }  
  
        context.write(key, new IntWritable(wordCount));  
    }  
}
```




The reduce Method: Processing the values

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
  
        int wordCount = 0;  
  
        for (IntWritable value : values) {  
            wordCount += value.get();  
        }  
  
        context.write(key, new IntWritable(wordCount));  
    }  
}
```

Step through all the elements in the collection. For the wordcounting example, just simply add all the values together.

The reduce method: Writing the Final Output

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
  
        int wordCount = 0;  
  
        for (IntWritable value : values) {  
            wordCount += value.get();  
        }  
  
        context.write(key, new IntWritable(wordCount));  
    }  
}
```



Write the final output
key-value pair to HDFS

The Driver: Complete Code

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.conf.Configuration;

public class WordCount {
    public static void main(String[] args) throws Exception {

        if (args.length != 2) {
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");
            System.exit(-1);
        }

        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "wordcount");
        job.setJarByClass(WordCount.class);
```

The Driver: Complete Code (Cont.)

```
FileInputFormat.setInputPaths(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);
```

```
job.setMapOutputKeyClass(Text.class);  
job.setMapOutputValueClass(IntWritable.class);
```

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);
```

```
}  
}
```


The Driver: Import Statements

```
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
import org.apache.hadoop.mapreduce.Job;  
Import org.apache.hadoop.conf.Configuration;
```

You will typically import the classes that are highlighted in red, in the driver. Other classes such as `IntWritable`, `Text` are imported if you need to use those classes in your program.

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "wordcount");  
        job.setJarByClass(WordCount.class);  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        boolean success = job.waitForCompletion(true);  
        System.exit(success ? 0 : 1);  
    }  
}
```

The *main* method accepts two command-line arguments: the input and output directories

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
    }  
}
```

Ensure two command-line arguments are provided (input dir, output dir)

```
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "wordcount");  
    job.setJarByClass(WordCount.class);  
  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    job.setMapperClass(WordMapper.class);  
    job.setReducerClass(SumReducer.class);  
  
    job.setMapOutputKeyClass(Text.class);  
    job.setMapOutputValueClass(IntWritable.class);  
  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
  
    boolean success = job.waitForCompletion(true);  
    System.exit(success ? 0 : 1);  
}
```

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
    }  
}
```

```
Configuration conf = new Configuration();  
Job job = Job.getInstance(conf, "wordcount");  
job.setJarByClass(WordCount.class);
```

```
FileInputFormat.setInputPaths(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);
```

```
job.setMapOutputKeyClass(Text.class);  
job.setMapOutputValueClass(IntWritable.class);
```

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);  
}  
}
```

- Create a new Job object
- Identify the Jar which contains the Mapper and Reducer by specifying a class in that jar
- Give the job a meaningful name

Creating a new Job object

- The Job and Configuration classes allow you to set configuration options for your MapReduce job
 - The classes to be used for your Mapper and Reducer
 - The input and output directories
 - Many other options
- Any options not explicitly set in your driver code will be read from your Hadoop configuration files
- Any options not specified in your configuration files will use Hadoop's default values
- You can also use the Job object to submit the job, control its execution, and query its state

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
    }  
}
```

```
Configuration conf = new Configuration();  
Job job = Job.getInstance(conf, "wordcount");  
job.setJarByClass(WordCount.class);
```

```
FileInputFormat.setInputPaths(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);
```

```
job.setMapOutputKeyClass(Text.class);  
job.setMapOutputValueClass(IntWritable.class);
```

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);  
}  
}
```

- Specify the input directory from which data will be read
 - All files in the directory are read
 - Exceptions: items whose names begin with a period (.) or underscore (_)
- Specify the output directory to which final output will be written
- Can specify any number of input files or directories as input
- The output path is a single directory

Configuring the Job: Specifying the InputFormat

- The default TextInputFormat will be used unless you specify otherwise
- TextInputFormat is the input format in which the value is a line of text and the key is the byte offset at which that line began in the file.
- To use an InputFormat other than the default
e.g.
`job.setInputFormatClass(KeyValueTextInputFormat.class)`
 - Until it finds the first delimiter (by default tab) it takes everything as key and remaining line as value

Configuring the Job: Specifying Final Output with OutputFormat

- The driver can specify the format of the output data
 - default is a plain text file (TextOutputFormat)
 - Can specify other format

e.g.

```
job.setOutputFormatClass(DBOutputFormat.class)
```

- Output to a DB table

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
    }  
}
```

```
Configuration conf = new Configuration();  
Job job = Job.getInstance(conf, "wordcount");  
job.setJarByClass(WordCount.class);  
  
FileInputFormat.setInputPaths(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);
```

```
job.setMapOutputKeyClass(Text.class);  
job.setMapOutputValueClass(IntWritable.class);
```

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);  
}  
}
```

- Give the Job object information about which class is Mapper, which class is Reducer
- Setting the Mapper and Reducer classes is optional
- If not set in the driver, Hadoop uses default
 - IdentityMapper
 - IdentityReducer

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
    }  
}
```

```
Configuration conf = new Configuration();  
Job job = Job.getInstance(conf, "wordcount");  
job.setJarByClass(WordCount.class);  
  
FileInputFormat.setInputPaths(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);
```

```
job.setMapOutputKeyClass(Text.class);  
job.setMapOutputValueClass(IntWritable.class);
```

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);  
}  
}
```

- Specify the types for the intermediate output keys and values produced by the Mapper.
- Must match the types of output keys and values in the class prototype of Mapper.
- If the types of the intermediate output key and value are identical to the reducer's output key and value (e.g. in this example), it is unnecessary to call `setMapOutputKeyClass` and `setMapOutputValueClass` methods.

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
    }  
}
```

```
Configuration conf = new Configuration();  
Job job = Job.getInstance(conf, "wordcount");  
job.setJarByClass(WordCount.class);
```

```
FileInputFormat.setInputPaths(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);
```

```
job.setMapOutputKeyClass(Text.class);  
job.setMapOutputValueClass(IntWritable.class);
```

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);  
}  
}
```

- Specify the types for the Reducer's output keys and values
- Must match the types of output keys and values in the class prototype of Reducer

The Driver: Main Code

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "wordcount");  
        job.setJarByClass(WordCount.class);  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        boolean success = job.waitForCompletion(true);  
        System.exit(success ? 0 : 1);  
    }  
}
```

- `waitForCompletion()`
 - Start the job and wait for it to complete (synchronous).
 - The parameter specifying verbosity: if true, display progress to the user.
 - The function polls the JobTracker for progress information and prints this to the console until the job is complete.
- `submit()`
 - does not block (driver code continues as the job is running)

In-Class Exercise

- Write a MapReduce job that reads any text input and computes the average length of all words that start with each character (case-sensitive).
- For example, for input

No now is definitely not the time

the output would be:

N	2.0
d	10.0
i	2.0
n	3.0
t	3.5

Discussion

- What are the key and value of input data to Mapper?
key: byte offset within the file; value: text of each line
- What are the key and value of output data of Mapper?
key: first letter of the word; value: length of the word
- What are the key and value of input data to Reducer?
key: first letter of the word;
value: a list of length of all words that start with the letter
- What are the key and value of output data of Reducer?
key: first letter of the word;
value: average length of all words that start with the letter

Now, let's write programs!