

# Midterm Notes

## Big Data Analytics - Midterm Exam Review

### 1. HDFS & Hadoop Cluster

#### HDFS Components

Hadoop Distributed File System (HDFS) is designed to store large amounts of data across multiple nodes in a distributed environment.

- **Data Node:** Stores actual file data in blocks.
  - Example: A 1GB file is split into 128MB blocks, each stored on different Data Nodes.
- **Name Node:** Manages metadata such as file names, locations, and block mappings.
  - Acts as the master for file system operations.
- **Replication:** Ensures fault tolerance by replicating data blocks across Data Nodes.
  - Default replication factor is **3**, meaning each block is stored on three different nodes.

#### Hadoop Cluster Architecture

A Hadoop cluster consists of different types of nodes:

- **Master Node:**
  - Runs the **Name Node**, which manages HDFS.
  - Runs the **Job Tracker**, which schedules jobs.
- **Slave Nodes:**
  - Run **Data Nodes**, which store data blocks.
  - Run **Task Trackers**, which execute MapReduce tasks.
- **Job Tracker:**
  - Assigns tasks to Task Trackers on Slave Nodes.
- **Task Tracker:**
  - Executes assigned tasks and reports status to the Job Tracker.

#### Storing and Retrieving Files

HDFS efficiently handles large files by splitting them into blocks.

##### 1. File Upload:

- A file is divided into **blocks** (e.g., 128MB each).
- Blocks are stored on **multiple Data Nodes** for redundancy.

##### 2. Read Operation:

- The client requests metadata from the **Name Node**.
- The client retrieves actual data directly from the **Data Nodes**.

#### Running a Job on a Hadoop Cluster

The process of executing a MapReduce job in Hadoop follows these steps:

1. The **client submits** a job to the **Job Tracker**.

2. The **Job Tracker** splits the job into **Map tasks** and **Reduce tasks**.
  3. **Task Trackers** execute tasks where data is stored, reducing network latency.
  4. The **Job Tracker** monitors task completion and manages failures.
  5. The **final output** is written back to **HDFS**.
- 

## 2. MapReduce

### MapReduce Workflow

MapReduce processes data using the following steps:

1. **Input Splitting**: The input file is divided into smaller chunks.
2. **Mapping**: Each Mapper processes its chunk and emits key-value pairs.
3. **Shuffling & Sorting**: The framework sorts intermediate data and assigns it to Reducers.
4. **Reducing**: Reducers aggregate results based on keys.
5. **Final Output**: The output is stored back in HDFS.

### Key Components

- **Mapper**:

- Processes input data line by line.
- Emits intermediate key-value pairs.
- **Example**: Word count

```
public class WordCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

- **Combiner** (Optional):

- Acts as a mini-reducer before data is sent to the final Reducer.
- Helps in reducing data transfer between Mapper and Reducer.

- **Partitioner**:

- Determines which Reducer receives a given key.
- **Example**: Assigning logs by month to different reducers.

- **Shuffle & Sort**:

- Intermediate key-value pairs are grouped by key before reaching the Reducer.

- **Reducer**:

- Aggregates values associated with each key.
- **Example**: Summing word counts

```

public class WordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

```

## Writing MapReduce Programs in Java

1. **Write a Mapper class:** `map()` method processes input and emits key-value pairs.
2. **Write a Reducer class:** `reduce()` method aggregates results based on keys.
3. **Write a Driver class:**
  - Sets the Mapper, Reducer, and Partitioner.
  - Specifies input and output paths.
  - Submits the job to Hadoop.

## MapReduce Flow & Combiners

- **Combiner:** Acts as a mini-reducer, performing local aggregation before sending data to the reducer.
  - Reduces data transfer between mapper and reducer.
  - Example: Word Count
    - Without Combiner: Each word count is sent individually.
    - With Combiner: Local aggregation occurs before sending to reducer.
- **Commutative & Associative Operations**
  - Commutative: Order of operation does not affect result (e.g., `a + b = b + a`).
  - Associative: Grouping of operation does not affect result (e.g., `(a + b) + c = a + (b + c)`).
  - **Examples:**
    - **Max function:** Can be used in a combiner (`max(max(A, B), max(C, D))`).
    - **Sum function:** Can also be used.
    - **Average function:** **Cannot** use a combiner directly because averaging partial values leads to incorrect results.

## Partitioners

- **Role of Partitioner:**
  - Determines which reducer receives a key.
  - **Default:** HashPartitioner distributes keys evenly.
  - **Custom Partitioner:** Used when specific grouping is needed (e.g., logs by month).
- **Examples of Partitioner Use Cases:**
  - Word count with multiple reducers.
  - Retail sales report with **12 reducers** (one per month).
  - Web traffic by day of the week with **7 reducers**.

- **Implementing a Custom Partitioner:**

- Extend `Partitioner<K, V>`.
- Override `getPartition()` method.
- Example:

```
public int getPartition(K key, V value, int numReduceTasks) {  
    return Math.abs(key.hashCode()) % numReduceTasks;  
}
```

- Can implement `Configurable` to set up variables before partitioning.

## Reducers

- **Single Reducer:**

- Advantage: Completely sorted output.
- Disadvantage: Slow if large data.

- **Multiple Reducers:**

- Example: Assign logs by month to 12 reducers.

---

## 3. Common MapReduce Algorithms

MapReduce is widely used for processing large-scale data in a distributed manner. Below are some common algorithms implemented using MapReduce, along with detailed explanations and examples.

### Sorting

Sorting is one of the most fundamental MapReduce operations. It is useful when dealing with large datasets where sorting helps in efficient querying and retrieval.

#### Example: Sorting Employee Salaries

**Input:**

```
Alice, 60000  
Bob, 50000  
Charlie, 70000  
David, 60000
```

#### Map Phase:

Each key-value pair is transformed to make salary the key so that sorting can be performed.

```
(60000, Alice)  
(50000, Bob)  
(70000, Charlie)  
(60000, David)
```

#### Shuffle & Sort Phase:

MapReduce automatically sorts the intermediate key-value pairs.

#### Reduce Phase:

The reducer outputs the sorted data in ascending order.

```
(50000, Bob)
(60000, Alice)
(60000, David)
(70000, Charlie)
```

Sorting can also be used in conjunction with secondary sorting when sorting within partitions is needed.

## Searching

Searching allows us to scan large text datasets and find relevant information efficiently.

### Example: Finding a Keyword in Log Files

**Input:** Log file contents

```
log1.txt: Error at line 45
log1.txt: Warning at line 100
log2.txt: Error at line 200
```

**Search pattern:** "Error"

#### Map Phase:

Each mapper scans a line and emits results if the search pattern is found.

```
(Error, log1.txt:45)
(Error, log2.txt:200)
```

#### Reduce Phase:

The reducer aggregates the results and outputs matching occurrences.

```
Error found in:
- log1.txt: Line 45
- log2.txt: Line 200
```

This approach is useful for searching error messages in logs, detecting spam in emails, or retrieving documents matching a keyword.

## Building an Inverted Index

An inverted index is a key data structure used in search engines to quickly find documents containing a given word.

### Example: Creating an Inverted Index for Documents

**Input:**

```
doc1.txt: Hadoop is a big data framework
doc2.txt: Hadoop uses MapRedu
```

#### Map Phase:

The mapper emits word-document pairs.

```
(Hadoop, doc1.txt)
(Hadoop, doc2.txt)
(big, doc1.txt)
```

```
(data, doc1.txt)
(framework, doc1.txt)
(uses, doc2.txt)
(MapReduce, doc2.txt)
```

### Reduce Phase:

The reducer groups words and outputs a list of documents where each word appears.

```
Hadoop → [doc1.txt, doc2.txt]
big → [doc1.txt]
data → [doc1.txt]
framework → [doc1.txt]
uses → [doc2.txt]
MapReduce → [doc2.txt]
```

This structure helps search engines like Google retrieve relevant documents based on keyword searches.

### Word Co-Occurrence

Word co-occurrence measures how often two words appear together in a dataset. This is useful for NLP applications, topic modeling, and recommendation systems.

### Example: Measuring Co-Occurrence in a Sentence

Input:

```
doc1.txt: data science is fun
doc2.txt: big data is useful
```

### Map Phase:

The mapper emits word pairs with a count of 1.

```
(data, science) → 1
(science, is) → 1
(is, fun) → 1
(big, data) → 1
(data, is) → 1
(is, useful) → 1
```

### Reduce Phase:

The reducer sums up the occurrences.

```
(data, science) → 1
(science, is) → 1
(is, fun) → 1
(big, data) → 1
(data, is) → 1
(is, useful) → 1
```

This helps in building word association models, such as for predicting the next word in a sentence.

### TF-IDF Computation

TF-IDF (Term Frequency-Inverse Document Frequency) is a key algorithm in text mining and information retrieval. It measures the importance of a word in a document relative to a collection of documents.

### Formula:

- **Term Frequency (TF):**  $TF = \frac{\text{Number of times term appears in a document}}{\text{Total words in the document}}$
- **Inverse Document Frequency (IDF):**  $IDF = \log\left(\frac{\text{Total number of documents}}{\text{Number of documents containing the term}}\right)$
- **TF-IDF Score:**  $TF-IDF = TF \times IDF$

### Example: Computing TF-IDF for "Hadoop"

#### Documents:

```
doc1.txt: Hadoop is a big data tool
doc2.txt: Hadoop and MapReduce work together
doc3.txt: Big data uses MapReduce
```

#### Step 1: Compute TF

```
TF(Hadoop, doc1) = 1/6
TF(Hadoop, doc2) = 1/5
TF(Hadoop, doc3) = 0
```

#### Step 2: Compute IDF

```
IDF(Hadoop) = log(3 / 2) = 0.176
```

#### Step 3: Compute TF-IDF

```
TF-IDF(Hadoop, doc1) = 1/6 * 0.176
TF-IDF(Hadoop, doc2) = 1/5 * 0.176
TF-IDF(Hadoop, doc3) = 0
```

The words with the highest TF-IDF values are considered more important for identifying the topic of each document.

### Secondary Sort

Secondary sorting allows sorting within a key's partition, often using a composite key.

### Example: Sorting People by Last Name and Birth Year

#### Input:

```
John Smith, 1985
Alice Brown, 1990
Bob Smith, 1983
Charlie Brown, 1988
```

### Map Phase:

The mapper outputs composite keys for sorting.

```
(Smith, 1985) → John  
(Smith, 1983) → Bob  
(Brown, 1990) → Alice  
(Brown, 1988) → Charli
```

### Partitioning and Sorting:

- **Partitioner** ensures all people with the same last name go to the same reducer.
- **Sort Comparator** sorts by birth year.

### Reduce Phase:

The reducer receives sorted values within each partition.

```
Brown:  
  - Charlie (1988)  
  - Alice (1990)  
Smith:  
  - Bob (1983)  
  - John (1985)
```

This approach is useful when sorting within a category, such as sorting transactions within a user's account.

---

These MapReduce algorithms form the backbone of big data processing in distributed systems, enabling tasks such as sorting, searching, indexing, and text analysis at scale.